

# Assignment 4

## Computer Architecture

Nidhi Hegde (180472)

### Part A:

#### LRU Policy

|         | No of L1 accesses | No of L2 accesses | No of L1 Cache Misses | No of L2 Cache Misses | Percentage of L2 DoF Blocks | Percentage of L2 Blocks (>= 2 hits) |
|---------|-------------------|-------------------|-----------------------|-----------------------|-----------------------------|-------------------------------------|
| perl    | 602217256         | 873034            | 873034                | 16704                 | 8.59076                     | 53.7751                             |
| bzip    | 701200790         | 8745365           | 8745365               | 4548977               | 67.2544                     | 52.2912                             |
| gcc     | 530755808         | 23224339          | 23224339              | 5035921               | 94.777                      | 49.6113                             |
| mcf     | 539049280         | 68828431          | 68828431              | 33576159              | 60.981                      | 18.1957                             |
| soplex  | 530837767         | 19552677          | 19552677              | 18814562              | 97.9148                     | 22.377                              |
| hmmer   | 669811655         | 3285938           | 3285938               | 1619858               | 95.3233                     | 29.6265                             |
| omnetpp | 591902829         | 14092639          | 14092639              | 10471631              | 80.9856                     | 36.7517                             |
| xalanc  | 569708840         | 16548039          | 16548039              | 2415618               | 49.7791                     | 61.8453                             |

Formula for calculating % of DoF:  $(\text{\#evicted entries with 0 hits} + \text{\# entries with 0 hits finally present in cache}) / \text{\# L2 blocks filled}$

We see that the LRU policy offers the lowest L1 miss rate for the Perl benchmark. The highest miss rate is seen in mcf benchmark. Perl shows the lowest percentage of dead-on-fill blocks. Xalanc shows the highest percentage of L2 blocks experiencing at least 2 hits before getting evicted.

### Part B:

#### SRRIP Policy

|         | No of L1 Cache Misses | No of L2 Cache Misses |
|---------|-----------------------|-----------------------|
| perl    | 873035                | 16705                 |
| bzip    | 8747090               | 4601535               |
| gcc     | 23223777              | 5026432               |
| mcf     | 68830487              | 34482580              |
| soplex  | 19548330              | 18759545              |
| hammer  | 3279062               | 1555564               |
| omnetpp | 14088666              | 10456956              |
| xalanc  | 16547131              | 2337919               |

We can see that this policy performs **better** than LRU Policy as it suffers a lesser number of L2 and L1 misses than LRU policy for most of the benchmarks.

REASON:

SRRIP policy addresses 2 of the problems – dead-on-fill misses and less frequently accessed blocks - encountered in the LRU Policy. Thus, we may obtain a better performance for this policy.

## NRU Policy

|         | No of L1 Cache Misses | No of L2 Cache Misses |
|---------|-----------------------|-----------------------|
| perl    | 873038                | 16724                 |
| bzip    | 8745671               | 4570825               |
| gcc     | 23224605              | 5039692               |
| mcf     | 68829771              | 33802251              |
| soplex  | 19553118              | 18818585              |
| hammer  | 3286045               | 1603119               |
| omnetpp | 14094205              | 10486398              |
| xalanc  | 16549366              | 2427413               |

This policy offers a slightly **worse** performance than LRU and SRRIP Policy on most of the benchmarks.

REASON:

We just maintain the most recently accessed way index in the NRU policy along with the REF bits. As soon as all REF bits are set to 1 (upon a hit or fill), we reset all of them except the most recently used entry. We do not differentiate between the second most recently used entry or the least recently used entry. If the former has a lower way index, we choose to evict it, though it may not be optimal. However, in case of LRU, we maintain additional information about the LRU state of all the entries, and only choose to evict the entry that was least recently accessed if need arises.

### Possible Reason for Different L1 misses:

The 3 cache replacement policies are different for L2 cache only. But the policies affect the number of misses in L1 because of the inclusive cache hierarchy. When a block is evicted from L2, its corresponding entry in L1 is also evicted. Since different replacement policies may choose to evict different blocks in L2, the L1 cache state is not conserved across different policies. Hence, when a particular block that was indirectly evicted from L1 by the policy in the past is requested again, it would add an extra L1 miss.